

# Day 1: Introduction to NumPy

## 1. What is NumPy?

**NumPy (Numerical Python)** is a powerful library for numerical computing in Python. It provides:

- A high-performance **n-dimensional array object** (`ndarray`)
- Tools for integrating with C/C++ and Fortran
- Functions for **linear algebra, Fourier transforms, random number generation**, and more

### **Why Data Engineers should learn NumPy:**

Data engineering often requires working with large datasets, doing fast matrix operations, and preparing data for ML. NumPy is **10x to 100x faster** than native Python lists due to vectorization and C-based backend.

---

## 2. Installing NumPy

You can install NumPy via pip or conda:

```
pip install numpy
```

Or, using Anaconda:

```
conda install numpy
```

---

## 3. Importing NumPy

The standard convention is:

```
import numpy as np
```

---

## 4. Why Not Just Use Python Lists?

Feature	Python List	NumPy Array
Type	Heterogeneous	Homogeneous
Speed	Slow	Very fast (compiled C backend)
Memory Usage	High	Efficient (contiguous memory)
Operations	Element-wise not easy	Vectorized and built-in

### Real-time Scenario:

When processing millions of machine sensor readings in real-time, NumPy arrays help you perform transformations (e.g., scaling, offset correction) with minimal delay.

---

## 5. Creating Your First NumPy Array

```
import numpy as np
```

```
arr = np.array([1, 2, 3, 4, 5])  
print(arr)
```

Output:

```
[1 2 3 4 5]
```

### Check Array Properties

```
print("Type:", type(arr))      # <class 'numpy.ndarray'>  
print("Shape:", arr.shape)    # (5,)   
print("Data Type:", arr.dtype) # int64 (or int32 depending on your system)
```

---

## 6. Data Engineer Use Case

**Scenario:** You have a machine producing 1 reading per second for 5 sensors.

```
sensor_readings = np.array([
    [20.5, 30.2, 40.1, 21.3, 25.5],
    [20.7, 30.1, 40.4, 21.5, 25.7],
    [20.9, 30.0, 40.6, 21.6, 25.9]
])

print("Shape:", sensor_readings.shape) # (3, 5)
```

This represents 3 seconds of readings for 5 sensors. You can now:

- Calculate averages
- Detect anomalies
- Reshape for time-based processing

---

## 7. Summary

Concept	Description
<code>np.array()</code>	Convert list to NumPy array
<code>.dtype</code>	Data type of array
<code>.shape</code>	Shape of the array (rows, columns)
<code>.ndim</code>	Number of dimensions
<code>type()</code>	Object type (should be ndarray)

---

## 8. Day 1 Task

1. Install NumPy and Jupyter Notebook (if not already done).

2. Convert the following list of hourly temperature readings into a NumPy array:

```
temps = [72.4, 73.0, 72.5, 74.2, 75.1, 76.5, 74.0]
```

3. Print:

- Mean temperature
- Minimum and maximum
- Total number of readings
- Data type and shape

4. BONUS: Create a 3x3 matrix of random machine speeds and print its shape and data type.

---

## 9. Real-Time Interview Tip

**Q:** Why would you prefer NumPy over Pandas for initial data cleaning in batch pipelines?

**A:** NumPy offers faster vectorized operations for raw numerical data. In the ETL layer, we often clean or normalize structured arrays before converting them into DataFrames for downstream analytics.